

```
In [4]: import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
import statsmodels.api as sm
from sklearn import linear_model
df = pd.read_csv(r'C:\Users\Berk\Desktop\apl_e.csv')
```

```
In [5]: df.tail(2)
```

Out[5]:

	#RIC	Date-Time	Price	Volume	Numberoftrades	Bidprice	Ask Price	Ni
24072	AAPL.OQ	2019-12-31T22:00:00.000000000Z	293.675	2716.0	NaN	293.50	293.85	
24073	AAPL.OQ	2019-12-31T23:00:00.000000000Z	293.810	848.0	NaN	293.77	293.85	

```
In [6]: df = df.dropna(subset=['Numberoftrades', 'Depth', 'Volume', 'Spread', 'Price'])
```

```
In [7]: x = df[['Numberoftrades', 'Depth', 'Volume', 'Spread']]
y = df[['Price']]
```

```
In [8]: x_normalized = StandardScaler().fit_transform(x)
y_normalized = StandardScaler().fit_transform(y)
```

```
In [9]: #To check if the data is normalized
x_normalized
```

```
Out[9]: array([[ -8.75524028e-01, -9.51987624e-01, -7.68765782e-01,
  1.91725535e-01],
 [ -8.79798866e-01, -9.52366910e-01, -7.70541302e-01,
  1.57982502e-01],
 [ -8.79086393e-01, -9.52499659e-01, -7.69804553e-01,
  7.76419457e-02],
 ...,
 [ -4.48277732e-01, -3.18827468e-01, -4.78221197e-01,
 -1.89520475e-03],
 [ 7.10203332e-01, 5.81216890e-01, 4.55488938e-01,
 -1.09179919e-03],
 [ -8.81223812e-01, -9.43870916e-01, 1.36051802e+00,
 1.17626898e-02]])
```

```
In [10]: x_e = pd.DataFrame(x_normalized, columns = ['Numberoftrades', 'Depth', 'Volume', 'Spread'])
```

```
In [11]: reg = linear_model.LinearRegression()  
reg.fit(x_normalized, y_normalized)
```

```
Out[11]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [12]: model = sm.OLS(y_normalized, x_normalized).fit()
         predictions = model.predict(x_normalized)

         print_model = model.summary()
         print(print_model)
```

# OLS Regression Results

```

=====
=====
Dep. Variable:          y      R-squared (uncentered):
0.144
Model:                  OLS    Adj. R-squared (uncentered):
0.144
Method:                 Least Squares    F-statistic:
673.1
Date:                   Wed, 03 Jun 2020    Prob (F-statistic):
0.00
Time:                   02:49:04    Log-Likelihood:
-21463.
No. Observations:      16003    AIC:
4.293e+04
Df Residuals:          15999    BIC:
4.296e+04
Df Model:              4
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
x1	0.3988	0.013	29.733	0.000	0.372	0.425
x2	-0.5490	0.012	-45.031	0.000	-0.573	-0.525
x3	-0.1995	0.009	-23.088	0.000	-0.216	-0.183
x4	-0.0151	0.007	-2.067	0.039	-0.029	-0.001

```

=====
Omnibus:               4614.893    Durbin-Watson:           0.107
Prob(Omnibus):         0.000    Jarque-Bera (JB):       10705.479
Skew:                  1.651    Prob(JB):                0.000
Kurtosis:              5.269    Cond. No.                3.46
=====

```

## Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [48]: from sklearn.tree import DecisionTreeRegressor
df = pd.read_csv(r'C:\Users\Berk\Desktop\apl_e.csv')
df.describe()
df = df.dropna(subset=['Numberoftrades', 'Depth', 'Volume', 'Spread', 'Price'])
x = df[['Numberoftrades', 'Depth', 'Volume', 'Spread']].values
y = df[['Price']].values
```

```
In [49]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)
regr = DecisionTreeRegressor()
regr.fit(x_train, y_train)
```

```
Out[49]: DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
                                max_leaf_nodes=None, min_impurity_decrease=0.0,
                                min_impurity_split=None, min_samples_leaf=1,
                                min_samples_split=2, min_weight_fraction_leaf=0.0,
                                presort=False, random_state=None, splitter='best')
```

```
In [50]: y_pred = regr.predict(x_test)
```

```
In [51]: df2=pd.DataFrame({'y_pred':y_pred})
df2.assign(y=y_test)
```

Out[51]:

	y_pred	y
0	518.10	528.75
1	223.64	186.70
2	112.55	114.95
3	265.98	215.61
4	98.20	175.32
...	...	...
3196	131.77	97.36
3197	110.67	120.25
3198	93.39	511.83
3199	131.50	114.61
3200	104.97	127.15

3201 rows × 2 columns

```
In [52]: from sklearn import metrics
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

Mean Absolute Error: 40.269442674164324  
Mean Squared Error: 6366.101425572007  
Root Mean Squared Error: 79.78785261912998

In [ ]: