

```
In [36]: import numpy as np
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
In [37]: df = pd.read_csv('https://raw.githubusercontent.com/carlson9/KocPython2020/master/homework/immSurvey.csv')
```

In [38]:

```
df
```

Out[38]:

Unnamed: 0	ids	X.1	textToSend	Means	stanMeans	X	MetalD	treatment	pid_rep	
0	1	6606	1	problems caused by the influx of illegal immig...	0.750000	2.409768	1	0	1	1.00000
1	2	6607	2	if you mean illegal immigration, i'm afraid of...	0.875000	3.710615	2	0	1	1.00000
2	3	6608	3	that they should enter the same way my grandpa...	0.416667	-1.437706	3	0	0	0.33300
3	4	6609	4	legally entering the usa meeting the requireme...	0.458333	0.655503	4	0	0	0.50000
4	5	6610	5	terror bombings killing us robbing america	0.875000	5.337525	5	0	1	0.66667
...	...	...	...	...	...	...	...	...	...	...
335	337	6939	337	those people that are here illegally should be...	0.650000	-0.326699	336	0	0	0.50000
336	338	6940	338	racism & xenophobia make me worried! also, co...	0.650000	0.789189	337	0	1	0.50000
337	339	6941	339	that immigrants are the heart and soul and the...	0.200000	-2.239257	338	0	0	0.50000
338	340	6942	340	job security for our citizens	0.300000	-1.716588	339	0	1	0.50000
339	341	6943	341	public safety, effect on the economy, whether ...	0.842105	4.268994	340	0	1	0.50000

340 rows × 38 columns

```
In [39]: alphas = df.stanMeansNewSysPooled
sample = df.textToSend

v = CountVectorizer()
x = v.fit_transform(sample)
```

```
In [40]: x
```

```
Out[40]: <340x1475 sparse matrix of type '<class 'numpy.int64'>'
with 6239 stored elements in Compressed Sparse Row format>
```

```
In [41]: pd.DataFrame(x.toarray(), columns=v.get_feature_names())
```

```
Out[41]:
```

	11	12	125	18	1b	600	95	able	abolition	aboration	...	wreckless	wrong	wrongly	›
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
335	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
336	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
337	0	0	0	0	0	0	0	0	0	0	...	0	1	0	
338	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
339	0	0	0	0	0	0	0	0	0	0	...	0	0	0	

340 rows × 1475 columns

```
In [42]: from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(x, alphas,
random_state=1)
```

```
In [43]: from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.gaussian_process.kernels import ConstantKernel, RBF

rbf = ConstantKernel(1.0) * RBF(length_scale=1.0)
gpr = GaussianProcessRegressor(kernel=rbf, alpha=1e-8)
gpr.fit(xtrain.toarray(), ytrain)
```

```
Out[43]: GaussianProcessRegressor(alpha=1e-08, copy_X_train=True,
kernel=1**2 * RBF(length_scale=1),
n_restarts_optimizer=0, normalize_y=False,
optimizer='fmin_l_bfgs_b', random_state=None)
```

```
In [44]: m, c = gpr.predict(xtest.toarray(), return_cov=True)
```

```
In [45]: print (m,c)
```

```
[-4.70368302e-01 -1.88451167e-01  1.79316802e-01 -1.08671571e+00
 7.72459599e-02  4.54517233e-03 -4.21716023e-02  5.49369072e-09
 2.98879892e-03  5.09086166e-02 -1.50141289e-01 -1.66857615e-02
 5.43721214e-01  1.23403909e-11 -6.22090133e-01 -4.13130805e-01
-7.64168129e-01  3.52373448e-01 -4.52882609e-01 -1.29262377e+00
-8.32969768e-02 -7.29553404e-01 -9.95216088e-01  3.68688097e-03
 8.54655530e-14  2.87784204e-02 -1.16016714e+00 -2.84864588e-01
-8.39988585e-02 -1.07503417e+00 -2.16794271e-04  4.15471760e-06
-1.09992427e-01  1.08702424e-04 -1.42257897e+00 -1.34403006e-02
-7.06904402e-02 -1.65489446e-01 -1.22983445e-03 -1.61471917e-03
-7.37694006e-01 -1.62698645e-01 -3.74296032e-01  3.02417965e-01
-8.35432977e-01  1.24294052e-01 -9.68070785e-01  5.50835571e-02
 5.53143826e-09 -1.87067406e-02  2.11023464e-02 -1.81546821e-01
-3.76305514e-02 -7.19630101e-02  9.31731214e-04 -1.26322946e+00
 1.22225891e+00 -5.69459234e-03 -2.12699263e-02 -2.32891069e-01
 2.21443661e-02  7.82059053e-06 -2.06860989e-01  1.88364332e-01
-1.88297873e-01 -1.03254425e-01 -8.24605644e-01  3.65161751e-04
-9.47848804e-02 -3.19408726e-01 -3.10947531e-03 -1.86026687e-02
-1.52503957e+00 -1.39136038e-01 -3.28390340e-03 -3.30650003e-01
-1.19234770e+00  7.33661635e-02 -3.47838643e-01 -5.62486050e-02
 2.90508862e-01 -4.67463932e-03  7.33291588e-04 -1.19234770e+00
 7.86244634e-02] [[ 4.79211237e-01 -5.96123032e-03 -2.26081337e-03 ... -1.55
211484e-05
 2.80687473e-10 -1.27233189e-04]
 [-5.96123032e-03  7.12749388e-01 -4.93387333e-03 ... -4.63993380e-06
 2.06233752e-10  1.84411957e-04]
 [-2.26081337e-03 -4.93387333e-03  3.80176056e-01 ...  3.69791304e-05
 7.16987580e-11  6.38047600e-04]
 ...
 [-1.55211484e-05 -4.63993380e-06  3.69791304e-05 ...  9.43235175e-01
-1.60513584e-12  5.52921394e-04]
 [ 2.80687251e-10  2.06233919e-10  7.16988691e-11 ... -1.60513578e-12
 9.99999950e-09 -6.18173568e-11]
 [-1.27233189e-04  1.84411957e-04  6.38047600e-04 ...  5.52921394e-04
-6.18173464e-11  9.38018398e-01]]]
```

```
In [54]: np.corrcoef(ytest,m)
```

```
Out[54]: array([[1.          , 0.63286061],
 [0.63286061, 1.          ]])
```

```
In [55]: #CORRELATION COEFFICIENT IS 0.6328
```

```
In [56]: bigram_vectorizer = CountVectorizer(ngram_range=(2, 2), token_pattern=r'\b\w+
\b', min_df=1)
```

```
In [57]: anl = bigram_vectorizer.build_analyzer()
```

```
In [58]: new_x = bigram_vectorizer.fit_transform(sample)
pd.DataFrame(new_x.toarray(), columns=bigram_vectorizer.get_feature_names())
```

Out[58]:

	1 3	1 became	1 boarders	1 difficult	1 immigrants	1 is	1 language	1 not	11 style	12 million	...	you come	yc c
0	0	0	0	0	0	0	0	0	0	0	...	0	
1	0	0	0	0	0	0	0	0	0	0	...	0	
2	0	0	0	0	0	0	0	0	0	0	...	0	
3	0	0	0	0	0	0	0	0	0	0	...	0	
4	0	0	0	0	0	0	0	0	0	0	...	0	
...	...	...	...	...	...	...	...	...	...	...	...	...	
335	0	0	0	0	0	0	0	0	0	0	...	0	
336	0	0	0	0	0	0	0	0	0	0	...	0	
337	0	0	0	0	0	0	0	0	0	0	...	0	
338	0	0	0	0	0	0	0	0	0	0	...	0	
339	0	0	0	0	0	0	0	0	0	0	...	0	

340 rows × 5209 columns

```
In [64]: xtrain, xtest, ytrain, ytest = train_test_split(new_x, alphas,
random_state=1)

rbf = ConstantKernel(1.0) * RBF(length_scale=1.0)
gpr = GaussianProcessRegressor(kernel=rbf, alpha=1e-8)
gpr.fit(xtrain.toarray(), ytrain)
```

Out[64]: GaussianProcessRegressor(alpha=1e-08, copy\_X\_train=True,  
kernel=1\*\*2 \* RBF(length\_scale=1),  
n\_restarts\_optimizer=0, normalize\_y=False,  
optimizer='fmin\_l\_bfgs\_b', random\_state=None)

```
In [65]: m, c = gpr.predict(xtest.toarray(), return_cov=True)
```

```
In [66]: np.corrcoef(ytest, m)
```

Out[66]: array([[1. , 0.45842178],  
[0.45842178, 1. ]])

```
In [67]: #BIAGRAM VECTORIZER REDUCED CORRELATION COEFFICIENT FROM 0.6328 to 0.4555
```

```
In [ ]:
```